

API Performance Dashboard Implementation

Dashboard Structure and Rationale

This dashboard is designed to rapidly identify and diagnose API performance issues by organizing information in a problem-oriented rather than tool-oriented way.

Key Dashboard Features

1. Service Health Overview

- **RED Metrics Summary:** Single-panel overview showing Request Rate, Error Rate, and Duration (p95) with color-coded health indicators
- **SLO Compliance:** Percentage of requests meeting latency and availability targets
- **Time Comparison:** Current metrics vs. same time yesterday/last week to identify anomalies

2. Request Performance Breakdown

- **Latency Heatmap:** Shows distribution of response times across endpoints and time
- **Endpoint Table:** Sortable table of endpoints with key metrics (volume, errors, latency)
- **Resource Correlation:** CPU/Memory/DB Connection usage overlaid with latency spikes
- **User Impact:** Affected users/tenants during performance degradation periods

3. Error Analysis

- **Error Timeline:** Error count by type and endpoint over time
- **Status Code Distribution:** Breakdown of HTTP status codes
- **Log Context:** Direct links to relevant error logs with context
- **Exception Patterns:** Common exception types and frequency

4. API Component Analysis

- **Latency Breakdown:** Visualization of where time is spent in the request lifecycle:
 - Database query time
 - External service call time
 - Application processing time
 - Network/serialization time
- **Bottleneck Identification:** Highlights the component contributing most to latency

5. Dependency Health

- **Database Performance:** Query execution time, connection usage, query volume
- **External Services:** Call volume, latency, and error rates for dependencies
- **Infrastructure:** Node-level metrics for relevant infrastructure components

6. Trace Explorer

- **Slowest Traces:** Automatically surfaces the slowest recent request traces
- **Error Traces:** Links to traces for recent error conditions
- **Trace Comparison:** Compare normal vs. slow request traces side-by-side

7. Debugging Accelerators

- **Common Issues Panel:** One-click filters for typical problems (DB slowness, auth issues)
- **Runbook Links:** Direct access to relevant troubleshooting guides
- **Historical Context:** Quick access to similar past incidents

Implementation Details

Metrics Configuration in Prometheus

- Custom API middleware to capture detailed timing at key request processing stages
- Histogram metrics for latency with appropriate bucket sizes
- Counter metrics for errors with detailed labels
- Gauge metrics for connection pools and resource utilization

Logs Configuration in Loki

- Structured logging format with consistent fields
- Log level filters to focus on actionable information
- Log queries pre-configured for common error patterns

Traces Configuration in Tempo

- Sampling strategy that ensures problematic requests are always traced
- Span naming conventions for consistent visualization
- Automatic instrumentation of database and HTTP clients

Alert Rules

- Multi-level alerting based on severity
- Rate of change alerts for early detection
- Compound alerts that combine multiple indicators

Expandability

- Template variables for environment, service, and endpoint selection
- Consistent label structure across all metrics
- Dashboard JSON can be parametrized and version-controlled
- Automated dashboard provisioning via Grafana API