

# Metrics, Logs, and Dashboard Approach

## Important Metrics for API Endpoint Performance

Following the RED Method (Rate, Errors, Duration) with additional context metrics:

### Primary Metrics

1. **Request Rate:** Requests per second, broken down by endpoint and method
  - *Why:* Indicates load and usage patterns; sudden changes may correspond to issues
2. **Error Rate:** Percentage of requests resulting in errors (4xx, 5xx)
  - *Why:* Directly impacts user experience; immediate indicator of service health
  - *Breakdown:* By status code, endpoint, and error type
3. **Duration (Latency):** Response time distributions with percentiles (p50, p90, p99)
  - *Why:* Shows performance degradation; high percentiles reveal edge case issues
  - *Components:* Break down by middleware, DB, external API calls

### Supporting Metrics

4. **Saturation Metrics:** Resource utilization connected to each endpoint
  - CPU, Memory, Connection pools, Thread utilization
  - *Why:* Helps identify resource constraints causing performance issues
5. **Dependency Metrics:** Performance of connected services
  - Database query times, external API latency
  - *Why:* Many API issues stem from downstream dependencies
6. **Business Context Metrics:** Payload size, user count, tenant information
  - *Why:* Adds context to technical metrics for faster debugging

## Log Strategy

### Log Selection Criteria

1. **Signal-to-noise ratio:** Keep logs that provide actionable information
2. **Unique context:** Prioritize logs that add information not available in metrics
3. **Error details:** Maintain detailed logs for error conditions
4. **Transaction boundaries:** Preserve logs marking start/end of key operations

### Logs to Add

1. **Structured Context:** Add JSON context to all logs with:
  - TraceID for correlation with traces
  - UserID/TenantID for request context
  - RequestID for grouping related logs
2. **Error Details:** Enhanced error logs with:
  - Stack traces for unexpected errors
  - Categorized error types for filtering
  - Relevant request parameters (sanitized)
3. **Performance Boundary Logs:** Mark entry/exit of critical operations with timing:
  - Database transactions
  - External service calls
  - Authentication/authorization steps

## Logs to Remove

1. **Debug noise:** Remove excessive debug logs in production
2. **Redundant information:** Eliminate logs that duplicate metrics
3. **Periodic status logs:** Replace with metrics for system status
4. **Sensitive data:** Remove or mask any PII or sensitive information

## Sample Dashboard Design

The API Performance Dashboard follows these design principles:

1. **Top-down approach:** Start with high-level health, drill down to components
2. **Correlation panels:** Place related metrics adjacently for easy comparison
3. **Context preservation:** Maintain selected time ranges when drilling down
4. **Actionable insights:** Include links to relevant logs and traces

## Dashboard Sections

## 1. **Overview Section**

- Service health summary (color-coded)
- Request volume and error trends
- SLO compliance indicators
- Top 5 slowest endpoints

## 2. **Endpoint Performance Section**

- Response time heatmap by endpoint
- Error rate by endpoint
- Resource utilization correlation
- Apdex score for user satisfaction

## 3. **Dependency Analysis Section**

- Database query performance
- External API call latency
- Cache hit/miss ratios
- Connection pool utilization

## 4. **Debug Acceleration Section**

- Recent error logs table
- Slow transaction traces
- Resource contention indicators
- One-click filters for common issues